

K-FoldDistilBert_Binary

March 25, 2026

```
[5]: import pandas as pd
import torch
from torch import nn
from torch.utils.data import DataLoader, Dataset
from transformers import DistilBertTokenizer, DistilBertModel, □
    ↪get_linear_schedule_with_warmup
from torch.optim import AdamW
from tqdm import tqdm
import numpy as np
from sklearn.model_selection import train_test_split, StratifiedKFold
from sklearn.metrics import (accuracy_score, classification_report, □
    ↪confusion_matrix,
                                roc_curve, auc, precision_recall_fscore_support)

import re
from bs4 import BeautifulSoup
import contractions
import matplotlib.pyplot as plt
import seaborn as sns
from imblearn.over_sampling import RandomOverSampler
from collections import Counter
from copy import deepcopy

# -----
# 1. Load & Inspect Dataset
# -----
try:
    df = pd.read_csv('annotated_dataset_for_binary_classifications.csv', □
        ↪encoding='utf-8')
except UnicodeDecodeError:
    df = pd.read_csv('annotated_dataset_for_binary_classifications.csv', □
        ↪encoding='latin1')

print(" Dataset loaded. Original columns:", df.columns.tolist())

df = df.rename(columns={'comment_Text': 'Review', 'Have issue': 'Have_issue'})

def convert_issue(issue):
```

```

        issue = str(issue).lower().strip()
        return 0 if issue in ['no', '0'] else 1

df['Have_issue'] = df['Have_issue'].apply(convert_issue)
df['Review'] = df['Review'].fillna('')

print("\nClass distribution (original):")
print(df['Have_issue'].value_counts())

plt.figure(figsize=(6, 4))
sns.countplot(data=df, x='Have_issue')
plt.title('Original Class Distribution')
plt.xticks([0, 1], ['No issue', 'Have_issue'])
plt.show()

# -----
# 2. Preprocessing
# -----
def preprocess_review(text):
    if not isinstance(text, str):
        text = str(text)
    text = contractions.fix(text)
    text = BeautifulSoup(text, "html.parser").get_text()
    text = re.sub(r'[^a-zA-Z0-9\s]', '', text)
    text = ' '.join(text.split())
    return text

df['ReviewP'] = df['Review'].apply(preprocess_review)
texts = df['ReviewP'].tolist()
labels = df['Have_issue'].tolist()

# -----
# 3. Train/Test Split
# -----
X_temp, X_test, y_temp, y_test = train_test_split(
    texts, labels, test_size=0.15, random_state=42, stratify=labels
)
X_train_raw, y_train_raw = np.array(X_temp), np.array(y_temp)

print(f"\n Dataset Splits:")
print(f"→ Full Train (for CV): {len(X_train_raw)} samples")
print(f"→ Held-Out Test: {len(X_test)} samples")
print(f"→ Train class (raw): {Counter(y_train_raw)} | Test: {Counter(y_test)}")

# -----
# 4. Dataset & Model
# -----

```

```

class TextClassificationDataset(Dataset):
    def __init__(self, texts, labels, tokenizer, max_length):
        self.texts = texts
        self.labels = labels
        self.tokenizer = tokenizer
        self.max_length = max_length

    def __len__(self):
        return len(self.texts)

    def __getitem__(self, idx):
        text = str(self.texts[idx])
        label = int(self.labels[idx])
        encoding = self.tokenizer(
            text, truncation=True, padding='max_length',
            max_length=self.max_length, return_tensors='pt'
        )
        return {
            'input_ids': encoding['input_ids'].flatten(),
            'attention_mask': encoding['attention_mask'].flatten(),
            'label': torch.tensor(label)
        }

# In model class:
class DistilBERTClassifier(nn.Module):
    def __init__(self, model_name, num_classes):
        super().__init__()
        self.distilbert = DistilBertModel.from_pretrained(model_name)
        self.dropout = nn.Dropout(0.1)
        self.fc = nn.Linear(self.distilbert.config.hidden_size, num_classes)

    def forward(self, input_ids, attention_mask):
        outputs = self.distilbert(input_ids=input_ids,
        ↪attention_mask=attention_mask)
        cls_output = outputs.last_hidden_state[:, 0]
        x = self.dropout(cls_output)
        logits = self.fc(x)
        return logits

# -----
# 5. Hyperparameters
# -----
model_name = 'distilbert-base-uncased' # Corrected to Hugging Face Hub model
num_classes = 2
max_length = 128
batch_size = 16
num_epochs = 6

```

```

learning_rate = 1e-5

device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
print(f"Using device: {device}")

tokenizer = DistilBertTokenizer.from_pretrained(model_name)

# -----
# 6. 5-Fold CV
# -----
print("\n Starting 5-Fold CV...")

k_folds = 5
skf = StratifiedKFold(n_splits=k_folds, shuffle=True, random_state=42)
cv_results = []
all_train_accuracies = []
all_val_accuracies = []

for fold, (train_idx, val_idx) in enumerate(skf.split(X_train_raw,
↳ y_train_raw), 1):
    print(f"\n{'='*50}\n  FOLD {fold}/{k_folds}\n{'='*50}")

    X_tr, X_val = X_train_raw[train_idx], X_train_raw[val_idx]
    y_tr, y_val = y_train_raw[train_idx], y_train_raw[val_idx]

    ros = RandomOverSampler(random_state=42)
    X_tr_res, y_tr_res = ros.fit_resample(X_tr.reshape(-1, 1), y_tr)
    X_tr_res, y_tr_res = X_tr_res.flatten().tolist(), y_tr_res.tolist()

    print(f"Fold {fold} → Train: {len(X_tr_res)} (oversampled), Val:
↳ {len(X_val)}")
    print(f"  → Train classes: {Counter(y_tr_res)} | Val: {Counter(y_val)}")

    train_dataset = TextClassificationDataset(X_tr_res, y_tr_res, tokenizer,
↳ max_length)
    val_dataset = TextClassificationDataset(X_val.tolist(), y_val.tolist(),
↳ tokenizer, max_length)
    train_loader = DataLoader(train_dataset, batch_size=batch_size,
↳ shuffle=True)
    val_loader = DataLoader(val_dataset, batch_size=batch_size, shuffle=False)

    model = DistilBERTClassifier(model_name, num_classes).to(device)
    optimizer = AdamW(model.parameters(), lr=learning_rate)
    total_steps = len(train_loader) * num_epochs
    scheduler = get_linear_schedule_with_warmup(optimizer, num_warmup_steps=0,
↳ num_training_steps=total_steps)

```

```

criterion = nn.CrossEntropyLoss()

fold_train_accs = []
fold_val_accs = []
best_val_acc = 0.0
best_state = None

for epoch in range(num_epochs):
    # Train
    model.train()
    total_loss, correct, total = 0, 0, 0
    for batch in train_loader:
        optimizer.zero_grad()
        input_ids = batch['input_ids'].to(device)
        attention_mask = batch['attention_mask'].to(device)
        labels = batch['label'].to(device)
        outputs = model(input_ids, attention_mask)
        loss = criterion(outputs, labels)
        loss.backward()
        torch.nn.utils.clip_grad_norm_(model.parameters(), 1.0)
        optimizer.step()
        scheduler.step()
        total_loss += loss.item()
        preds = outputs.argmax(dim=1)
        correct += (preds == labels).sum().item()
        total += labels.size(0)
    train_acc = correct / total
    fold_train_accs.append(train_acc)

    # Validate
    model.eval()
    val_preds, val_true = [], []
    with torch.no_grad():
        for batch in val_loader:
            input_ids = batch['input_ids'].to(device)
            attention_mask = batch['attention_mask'].to(device)
            labels = batch['label'].to(device)
            outputs = model(input_ids, attention_mask)
            preds = outputs.argmax(dim=1)
            val_preds.extend(preds.cpu().numpy())
            val_true.extend(labels.cpu().numpy())
    val_acc = accuracy_score(val_true, val_preds)
    fold_val_accs.append(val_acc)

    if val_acc > best_val_acc:
        best_val_acc = val_acc
        best_state = deepcopy(model.state_dict())

```

```

        print(f"Epoch {epoch+1}/{num_epochs} | Train Acc: {train_acc:.4f} | Val_
↪Acc: {val_acc:.4f}")

    all_train_accuracies.append(fold_train_accs)
    all_val_accuracies.append(fold_val_accs)

    model.load_state_dict(best_state)
    precision, recall, f1, _ = precision_recall_fscore_support(val_true,
↪val_preds, average='binary')
    cv_results.append({'fold': fold, 'val_accuracy': best_val_acc, 'precision':
↪precision, 'recall': recall, 'f1': f1})

# CV Summary
cv_df = pd.DataFrame(cv_results)
print("\n" + "="*60)
print(" K-FOLD CV RESULTS (mean ± std):")
print

```

Dataset loaded. Original columns: ['Rating_Star', 'Headings', 'comment_Text', 'Have_issue']


Class distribution (original):

Have_issue

0 45096

1 16198

Name: count, dtype: int64



Dataset Splits:

→ Full Train (for CV): 52099 samples
 → Held-Out Test: 9195 samples
 → Train class (raw): Counter({0: 38331, 1: 13768}) | Test: Counter({0: 6765, 1: 2430})
 Using device: cpu

Starting 5-Fold CV...

```
=====
FOLD 1/5
=====
Fold 1 → Train: 61328 (oversampled), Val: 10420
  → Train classes: Counter({0: 30664, 1: 30664}) | Val: Counter({0: 7667, 1: 2753})
Epoch 1/6 | Train Acc: 0.9536 | Val Acc: 0.9735
Epoch 2/6 | Train Acc: 0.9887 | Val Acc: 0.9742
Epoch 3/6 | Train Acc: 0.9944 | Val Acc: 0.9775
Epoch 4/6 | Train Acc: 0.9971 | Val Acc: 0.9789
Epoch 5/6 | Train Acc: 0.9984 | Val Acc: 0.9789
Epoch 6/6 | Train Acc: 0.9989 | Val Acc: 0.9794
```

```

=====
FOLD 2/5
=====
Fold 2 → Train: 61330 (oversampled), Val: 10420
  → Train classes: Counter({0: 30665, 1: 30665}) | Val: Counter({0: 7666, 1:
2754})
Epoch 1/6 | Train Acc: 0.9541 | Val Acc: 0.9713
Epoch 2/6 | Train Acc: 0.9878 | Val Acc: 0.9765
Epoch 3/6 | Train Acc: 0.9948 | Val Acc: 0.9778
Epoch 4/6 | Train Acc: 0.9967 | Val Acc: 0.9799
Epoch 5/6 | Train Acc: 0.9985 | Val Acc: 0.9790
Epoch 6/6 | Train Acc: 0.9989 | Val Acc: 0.9789

=====
FOLD 3/5
=====
Fold 3 → Train: 61330 (oversampled), Val: 10420
  → Train classes: Counter({1: 30665, 0: 30665}) | Val: Counter({0: 7666, 1:
2754})
Epoch 1/6 | Train Acc: 0.9542 | Val Acc: 0.9671
Epoch 2/6 | Train Acc: 0.9874 | Val Acc: 0.9748
Epoch 3/6 | Train Acc: 0.9949 | Val Acc: 0.9775
Epoch 4/6 | Train Acc: 0.9975 | Val Acc: 0.9797
Epoch 5/6 | Train Acc: 0.9986 | Val Acc: 0.9767
Epoch 6/6 | Train Acc: 0.9991 | Val Acc: 0.9782

=====
FOLD 4/5
=====
Fold 4 → Train: 61330 (oversampled), Val: 10420
  → Train classes: Counter({0: 30665, 1: 30665}) | Val: Counter({0: 7666, 1:
2754})
Epoch 1/6 | Train Acc: 0.9546 | Val Acc: 0.9630
Epoch 2/6 | Train Acc: 0.9880 | Val Acc: 0.9721
Epoch 3/6 | Train Acc: 0.9947 | Val Acc: 0.9759
Epoch 4/6 | Train Acc: 0.9971 | Val Acc: 0.9759
Epoch 5/6 | Train Acc: 0.9982 | Val Acc: 0.9798
Epoch 6/6 | Train Acc: 0.9990 | Val Acc: 0.9792

=====
FOLD 5/5
=====
Fold 5 → Train: 61330 (oversampled), Val: 10419
  → Train classes: Counter({0: 30665, 1: 30665}) | Val: Counter({0: 7666, 1:
2753})
Epoch 1/6 | Train Acc: 0.9559 | Val Acc: 0.9685
Epoch 2/6 | Train Acc: 0.9878 | Val Acc: 0.9752
Epoch 3/6 | Train Acc: 0.9944 | Val Acc: 0.9766

```


Epoch 4/6 | Train Acc: 0.9972 | Val Acc: 0.9801
 Epoch 5/6 | Train Acc: 0.9983 | Val Acc: 0.9809
 Epoch 6/6 | Train Acc: 0.9990 | Val Acc: 0.9808

=====

K-FOLD CV RESULTS (mean ± std):

[5]: <function print(*args, sep=' ', end='\n', file=None, flush=False)>

```
[6]: # -----
# 7. Final Model Training with Internal Val
# -----
print("\n Retraining final model on FULL TRAIN...")

ros_full = RandomOverSampler(random_state=42)
X_train_res, y_train_res = ros_full.fit_resample(np.array(X_train_raw).
    ↳reshape(-1, 1), y_train_raw)
X_train_res, y_train_res = X_train_res.flatten().tolist(), y_train_res.tolist()

# Internal val split
X_tr_final, X_val_final, y_tr_final, y_val_final = train_test_split(
    X_train_res, y_train_res, test_size=0.1, random_state=42,
    ↳stratify=y_train_res
)

train_dataset = TextClassificationDataset(X_tr_final, y_tr_final, tokenizer,
    ↳max_length)
val_dataset = TextClassificationDataset(X_val_final, y_val_final, tokenizer,
    ↳max_length)
test_dataset = TextClassificationDataset(X_test, y_test, tokenizer, max_length)

train_loader = DataLoader(train_dataset, batch_size=batch_size, shuffle=True)
val_loader = DataLoader(val_dataset, batch_size=batch_size, shuffle=False)
test_loader = DataLoader(test_dataset, batch_size=batch_size, shuffle=False)

final_model = DistilBERTClassifier(model_name, num_classes).to(device)
optimizer = AdamW(final_model.parameters(), lr=learning_rate)
total_steps = len(train_loader) * num_epochs
scheduler = get_linear_schedule_with_warmup(optimizer, num_warmup_steps=0,
    ↳num_training_steps=total_steps)
criterion = nn.CrossEntropyLoss()

# Track final history
final_train_losses, final_val_losses = [], []
final_train_accs, final_val_accs = [], []

for epoch in range(num_epochs):
```

```

# Train
final_model.train()
total_loss, correct, total = 0, 0, 0
for batch in tqdm(train_loader, desc=f"Epoch {epoch+1}/{num_epochs}"):
    optimizer.zero_grad()
    input_ids = batch['input_ids'].to(device)
    attention_mask = batch['attention_mask'].to(device)
    labels = batch['label'].to(device)
    outputs = final_model(input_ids, attention_mask)
    loss = criterion(outputs, labels)
    loss.backward()
    torch.nn.utils.clip_grad_norm_(final_model.parameters(), 1.0)
    optimizer.step()
    scheduler.step()
    total_loss += loss.item()
    preds = outputs.argmax(dim=1)
    correct += (preds == labels).sum().item()
    total += labels.size(0)
avg_train_loss = total_loss / len(train_loader)
train_acc = correct / total
final_train_losses.append(avg_train_loss)
final_train_accs.append(train_acc)

# Validate
final_model.eval()
val_loss, val_correct, val_total = 0, 0, 0
with torch.no_grad():
    for batch in val_loader:
        input_ids = batch['input_ids'].to(device)
        attention_mask = batch['attention_mask'].to(device)
        labels = batch['label'].to(device)
        outputs = final_model(input_ids, attention_mask)
        loss = criterion(outputs, labels)
        val_loss += loss.item()
        preds = outputs.argmax(dim=1)
        val_correct += (preds == labels).sum().item()
        val_total += labels.size(0)
avg_val_loss = val_loss / len(val_loader)
val_acc = val_correct / val_total
final_val_losses.append(avg_val_loss)
final_val_accs.append(val_acc)

# -----
# 7b. Final Model Learning Curves
# -----
epochs = range(1, num_epochs + 1)
plt.figure(figsize=(14, 5))

```

```

plt.subplot(1, 2, 1)
plt.plot(epochs, final_train_accs, 'b-o', label='Train Acc')
plt.plot(epochs, final_val_accs, 'r-o', label='Val Acc')
plt.title('Final DistilBERT Model: Accuracy')
plt.xlabel('Epoch'); plt.ylabel('Accuracy'); plt.legend(); plt.grid(True)

plt.subplot(1, 2, 2)
plt.plot(epochs, final_train_losses, 'b-o', label='Train Loss')
plt.plot(epochs, final_val_losses, 'r-o', label='Val Loss')
plt.title('Final DistilBERT Model: Loss')
plt.xlabel('Epoch'); plt.ylabel('Loss'); plt.legend(); plt.grid(True)

plt.tight_layout()
plt.savefig('distilbertBinary_final_learning_curves.png', dpi=300,
    ↳bbox_inches='tight')
plt.show()

# -----
# 8. Final Test Evaluation
# -----
print("\n Evaluating on HELD-OUT TEST SET...")
final_model.eval()
test_preds, test_labels = [], []
with torch.no_grad():
    for batch in tqdm(test_loader, desc="Test Evaluation"):
        input_ids = batch['input_ids'].to(device)
        attention_mask = batch['attention_mask'].to(device)
        labels = batch['label'].to(device)
        outputs = final_model(input_ids, attention_mask)
        preds = outputs.argmax(dim=1)
        test_preds.extend(preds.cpu().numpy())
        test_labels.extend(labels.cpu().numpy())

test_acc = accuracy_score(test_labels, test_preds)
precision, recall, f1, _ = precision_recall_fscore_support(test_labels,
    ↳test_preds, average='binary')

print(f"\n FINAL TEST RESULTS:")
print(f"Accuracy : {test_acc:.4f}")
print(f"Precision: {precision:.4f}")
print(f"Recall : {recall:.4f}")
print(f"F1-Score : {f1:.4f}")

print("\n Classification Report (Test Set):")
print(classification_report(test_labels, test_preds, target_names=["No issue",
    ↳"Have_issue"]))

```

```

# Confusion Matrix
cm = confusion_matrix(test_labels, test_preds)
plt.figure(figsize=(6, 5))
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues', xticklabels=["No issue", "Have issue"], yticklabels=["No issue", "Have issue"])
plt.title("DistilBERT: Confusion Matrix (Final Test Set)")
plt.ylabel("True Label"); plt.xlabel("Predicted Label")
plt.tight_layout()
plt.savefig('distilbertbinary_final_cm.png', dpi=300, bbox_inches='tight')
plt.show()

# ROC Curve
fpr, tpr, _ = roc_curve(test_labels, test_preds)
roc_auc = auc(fpr, tpr)
plt.figure(figsize=(6, 5))
plt.plot(fpr, tpr, color='darkorange', lw=2, label=f'ROC (AUC = {roc_auc:.3f})')
plt.plot([0, 1], [0, 1], color='navy', lw=2, linestyle='--')
plt.xlim([0.0, 1.0]); plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate'); plt.ylabel('True Positive Rate')
plt.title('DistilBERT: ROC Curve (Final Test Set)')
plt.legend(loc="lower right"); plt.grid(True)
plt.tight_layout()
plt.savefig('distilbertBinary_final_roc.png', dpi=300, bbox_inches='tight')
plt.show()

```

Retraining final model on FULL TRAIN...

Epoch 1/6: 100%	4313/4313 [1:31:43<00:00, 1.28s/it]
Epoch 2/6: 100%	4313/4313 [1:26:29<00:00, 1.20s/it]
Epoch 3/6: 100%	4313/4313 [1:25:58<00:00, 1.20s/it]
Epoch 4/6: 100%	4313/4313 [1:26:00<00:00, 1.20s/it]
Epoch 5/6: 100%	4313/4313 [1:26:19<00:00, 1.20s/it]
Epoch 6/6: 100%	4313/4313 [1:26:11<00:00, 1.20s/it]

Evaluating on HELD-OUT TEST SET...

Test Evaluation: 100% | 575/575 [02:56<00:00, 3.26it/s]

FINAL TEST RESULTS:

Accuracy : 0.9822

Precision: 0.9521

Recall : 0.9819

F1-Score : 0.9668

Classification Report (Test Set):

	precision	recall	f1-score	support
No issue	0.99	0.98	0.99	6765
Have_issue	0.95	0.98	0.97	2430
accuracy			0.98	9195
macro avg	0.97	0.98	0.98	9195
weighted avg	0.98	0.98	0.98	9195

