

```

!pip install contractions

print("Hello World")

Hello World

## Load all libraries
import pandas as pd
import torch
from torch import nn
from torch.utils.data import Dataset, DataLoader
from transformers import RobertaTokenizer, RobertaModel, AdamW,
get_linear_schedule_with_warmup
import numpy as np
from sklearn.model_selection import KFold
from sklearn.metrics import accuracy_score, classification_report,
confusion_matrix, roc_curve, auc
import re
from bs4 import BeautifulSoup
import contractions
from sklearn.utils.class_weight import compute_class_weight
from imblearn.over_sampling import RandomOverSampler
import matplotlib.pyplot as plt
import seaborn as sns
import warnings

# Suppress warnings
warnings.filterwarnings("ignore")

# Load the dataset
file_path =
"/root/workspace/aka_project/Naikdil/Datset_for_binary.csv" # Update
path as needed
df = pd.read_csv(file_path)

# Check and rename columns
print("Original columns:", df.columns)
df = df.rename(columns={'Base_Reviews': 'Review', 'Have_issue':
'Issue'})

# Convert issues to binary (1 for any issue, 0 for no issue)
def convert_issue(issue):
    issue = str(issue).lower().strip()
    if issue == 'no' or issue == '0':
        return 0
    return 1 # Treat everything else as an issue

df['Issue'] = df['Issue'].apply(convert_issue)

# Check class distribution
class_counts = df['Issue'].value_counts()

```

```

print("\nClass Distribution:")
print(class_counts)

# Plot distribution
plt.figure(figsize=(8, 4))
sns.countplot(data=df, x='Issue')
plt.title('Distribution of Issues (1=Issue, 0=No Issue)')
plt.xticks([0, 1], ['No Issue', 'Issue'])
plt.show()

### Text Preprocessing
def preprocess_review(text):
    text = str(text)
    # Fix contractions
    text = contractions.fix(text)
    # Remove HTML
    text = BeautifulSoup(text, "html.parser").get_text()
    # Clean special characters
    text = re.sub(r'[^a-zA-Z0-9\s]', ' ', text)
    # Normalize whitespace
    text = re.sub(r'\s+', ' ', text).strip()
    return text

df['Clean_Review'] = df['Review'].apply(preprocess_review)

### Dataset Class
class TextClassificationDataset(Dataset):
    def __init__(self, texts, labels, tokenizer, max_length):
        self.texts = texts
        self.labels = labels
        self.tokenizer = tokenizer
        self.max_length = max_length

    def __len__(self):
        return len(self.texts)

    def __getitem__(self, idx):
        text = str(self.texts[idx])
        label = int(self.labels[idx])
        encoding = self.tokenizer(
            text,
            max_length=self.max_length,
            padding='max_length',
            truncation=True,
            return_tensors='pt'
        )
        return {
            'input_ids': encoding['input_ids'].flatten(),
            'attention_mask': encoding['attention_mask'].flatten(),
            'label': torch.tensor(label, dtype=torch.long)
        }

```

```

    }

### Model Architecture
class RoBERTaClassifier(nn.Module):
    def __init__(self, roberta_model_name, num_classes):
        super(RoBERTaClassifier, self).__init__()
        self.roberta =
        RobertaModel.from_pretrained(roberta_model_name)
        self.dropout = nn.Dropout(0.1)
        self.fc = nn.Linear(self.roberta.config.hidden_size,
num_classes)

    def forward(self, input_ids, attention_mask):
        outputs = self.roberta(input_ids=input_ids,
attention_mask=attention_mask)
        pooled_output = outputs.pooler_output
        x = self.dropout(pooled_output)
        logits = self.fc(x)
        return logits

### Training Function
def train(model, data_loader, optimizer, scheduler, device,
class_weights=None):
    model.train()
    criterion = nn.CrossEntropyLoss(weight=torch.tensor(class_weights,
dtype=torch.float).to(device))
    total_loss = 0.0
    total_correct = 0
    total_samples = 0

    for batch in data_loader:
        optimizer.zero_grad()
        input_ids = batch['input_ids'].to(device)
        attention_mask = batch['attention_mask'].to(device)
        labels = batch['label'].to(device)

        outputs = model(input_ids=input_ids,
attention_mask=attention_mask)
        loss = criterion(outputs, labels)
        loss.backward()
        torch.nn.utils.clip_grad_norm_(model.parameters(), 1.0)
        optimizer.step()
        scheduler.step()

        _, predicted = torch.max(outputs, 1)
        total_correct += (predicted == labels).sum().item()
        total_samples += labels.size(0)

    avg_loss = total_loss / len(data_loader)
    avg_accuracy = total_correct / total_samples

```

```

        return avg_loss, avg_accuracy

### Evaluation Function
def evaluate(model, data_loader, device):
    model.eval()
    predictions = []
    actual_labels = []
    total_loss = 0.0

    with torch.no_grad():
        for batch in data_loader:
            input_ids = batch['input_ids'].to(device)
            attention_mask = batch['attention_mask'].to(device)
            labels = batch['label'].to(device)

            outputs = model(input_ids=input_ids,
                             attention_mask=attention_mask)
            _, preds = torch.max(outputs, dim=1)
            predictions.extend(preds.cpu().tolist())
            actual_labels.extend(labels.cpu().tolist())
            loss = nn.CrossEntropyLoss()(outputs, labels)
            total_loss += loss.item()

    accuracy = accuracy_score(actual_labels, predictions)
    report = classification_report(actual_labels, predictions,
                                   target_names=['No Issue', 'Issue'], output_dict=True)
    avg_loss = total_loss / len(data_loader)
    return accuracy, report, avg_loss, predictions, actual_labels

### Main Training Setup
# Parameters
roberta_model_name = "roberta-base"
num_classes = 2
max_length = 128
batch_size = 16
num_epochs = 6
learning_rate = 1e-5

# Initialize
tokenizer = RobertaTokenizer.from_pretrained(roberta_model_name)
device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
model = RoBERTaClassifier(roberta_model_name, num_classes).to(device)
optimizer = AdamW(model.parameters(), lr=learning_rate)

# Prepare data
texts = df['Clean_Review'].values
labels = df['Issue'].values

# Compute class weights
class_weights = compute_class_weight('balanced',

```

```

classes=np.unique(labels), y=labels)
print("\nClass Weights:", class_weights)

# Cross-validation
n_splits = 5
kf = KFold(n_splits=n_splits, random_state=42, shuffle=True)

# Track metrics
all_predictions = []
all_actual_labels = []

for fold, (train_idx, val_idx) in enumerate(kf.split(texts)):
    print(f"\n=== Fold {fold + 1}/{n_splits} ===")

    # Split data
    train_texts, val_texts = texts[train_idx], texts[val_idx]
    train_labels, val_labels = labels[train_idx], labels[val_idx]

    # Handle class imbalance
    ros = RandomOverSampler(random_state=42)
    train_texts_resampled = train_texts.reshape(-1, 1)
    train_texts_resampled, train_labels_resampled =
ros.fit_resample(train_texts_resampled, train_labels)
    train_texts = train_texts_resampled[:, 0]
    train_labels = train_labels_resampled

    # Create datasets
    train_dataset = TextClassificationDataset(train_texts,
train_labels, tokenizer, max_length)
    val_dataset = TextClassificationDataset(val_texts, val_labels,
tokenizer, max_length)

    # Create dataloaders
    train_dataloader = DataLoader(train_dataset,
batch_size=batch_size, shuffle=True)
    val_dataloader = DataLoader(val_dataset, batch_size=batch_size)

    # Scheduler
    total_steps = len(train_dataloader) * num_epochs
    scheduler = get_linear_schedule_with_warmup(optimizer,
num_warmup_steps=0, num_training_steps=total_steps)

    # Training loop
    for epoch in range(num_epochs):
        print(f"\nEpoch {epoch + 1}/{num_epochs}")
        train_loss, train_accuracy = train(model, train_dataloader,
optimizer, scheduler, device, class_weights)
        print(f"Train Loss: {train_loss:.4f}, Accuracy:
{train_accuracy:.4f}")

```

```

        val_accuracy, val_report, val_loss, val_preds, val_true =
evaluate(model, val_dataloader, device)
        print(f"Val Loss: {val_loss:.4f}, Accuracy:
{val_accuracy:.4f}")
        print(classification_report(val_true, val_preds,
target_names=['No Issue', 'Issue']))

        all_predictions.extend(val_preds)
        all_actual_labels.extend(val_true)

### Final Evaluation
# Confusion Matrix
cm = confusion_matrix(all_actual_labels, all_predictions)
plt.figure(figsize=(6, 6))
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues',
            xticklabels=['No Issue', 'Issue'],
            yticklabels=['No Issue', 'Issue'])
plt.title('Confusion Matrix')
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.show()

# ROC Curve
fpr, tpr, _ = roc_curve(all_actual_labels, all_predictions)
roc_auc = auc(fpr, tpr)

plt.figure()
plt.plot(fpr, tpr, color='darkorange', lw=2, label=f'ROC curve (AUC =
{roc_auc:.2f})')
plt.plot([0, 1], [0, 1], color='navy', lw=2, linestyle='--')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('ROC Curve')
plt.legend(loc="lower right")
plt.show()


# Classification Report
print("\n=== Final Classification Report ===")
print(classification_report(all_actual_labels, all_predictions,
target_names=['No Issue', 'Issue']))

Original columns: Index(['Profile_Name', 'Rating_Star', 'Headings',
'Issue_D', 'Base_Reviews',
'Chagpt annoations', 'category', 'Final annoations', 'Issue
Details',
'Have_issue'],
dtype='object')

Class Distribution:
Issue

```

```
0    45096
1     16198
Name: count, dtype: int64
```



Some weights of RobertaModel were not initialized from the model checkpoint at roberta-base and are newly initialized:
['roberta.pooler.dense.bias', 'roberta.pooler.dense.weight']
You should probably TRAIN this model on a down-stream task to be able to use it for predictions and inference.

Class Weights: [0.67959464 1.89202371]

=== Fold 1/5 ===

Epoch 1/6

Train Loss: 0.0000, Accuracy: 0.9561

Val Loss: 0.1254, Accuracy: 0.9713

	precision	recall	f1-score	support
No Issue	0.99	0.97	0.98	9035
Issue	0.92	0.97	0.95	3224
accuracy			0.97	12259
macro avg	0.96	0.97	0.96	12259
weighted avg	0.97	0.97	0.97	12259

Epoch 2/6

Train Loss: 0.0000, Accuracy: 0.9866

Val Loss: 0.1236, Accuracy: 0.9786

	precision	recall	f1-score	support
No Issue	0.99	0.98	0.99	9035
Issue	0.94	0.98	0.96	3224
accuracy			0.98	12259
macro avg	0.97	0.98	0.97	12259
weighted avg	0.98	0.98	0.98	12259

Epoch 3/6

Train Loss: 0.0000, Accuracy: 0.9927

Val Loss: 0.1362, Accuracy: 0.9784

	precision	recall	f1-score	support
No Issue	0.99	0.98	0.99	9035
Issue	0.94	0.98	0.96	3224
accuracy			0.98	12259
macro avg	0.97	0.98	0.97	12259
weighted avg	0.98	0.98	0.98	12259

Epoch 4/6

Train Loss: 0.0000, Accuracy: 0.9962

Val Loss: 0.1397, Accuracy: 0.9806

	precision	recall	f1-score	support
No Issue	0.99	0.98	0.99	9035
Issue	0.94	0.98	0.96	3224
accuracy			0.98	12259
macro avg	0.97	0.98	0.98	12259
weighted avg	0.98	0.98	0.98	12259

Epoch 5/6

Train Loss: 0.0000, Accuracy: 0.9977

Val Loss: 0.1393, Accuracy: 0.9835

	precision	recall	f1-score	support
No Issue	0.99	0.99	0.99	9035
Issue	0.96	0.98	0.97	3224
accuracy			0.98	12259
macro avg	0.98	0.98	0.98	12259
weighted avg	0.98	0.98	0.98	12259

Epoch 6/6

Train Loss: 0.0000, Accuracy: 0.9986

Val Loss: 0.1439, Accuracy: 0.9834

	precision	recall	f1-score	support
No Issue	0.99	0.98	0.99	9035
Issue	0.96	0.98	0.97	3224
accuracy			0.98	12259
macro avg	0.98	0.98	0.98	12259
weighted avg	0.98	0.98	0.98	12259

=== Fold 2/5 ===

Epoch 1/6

Train Loss: 0.0000, Accuracy: 0.9916

Val Loss: 0.0134, Accuracy: 0.9971

	precision	recall	f1-score	support
No Issue	1.00	1.00	1.00	8953
Issue	0.99	1.00	0.99	3306
accuracy			1.00	12259
macro avg	0.99	1.00	1.00	12259
weighted avg	1.00	1.00	1.00	12259

Epoch 2/6

Train Loss: 0.0000, Accuracy: 0.9961

Val Loss: 0.0250, Accuracy: 0.9957

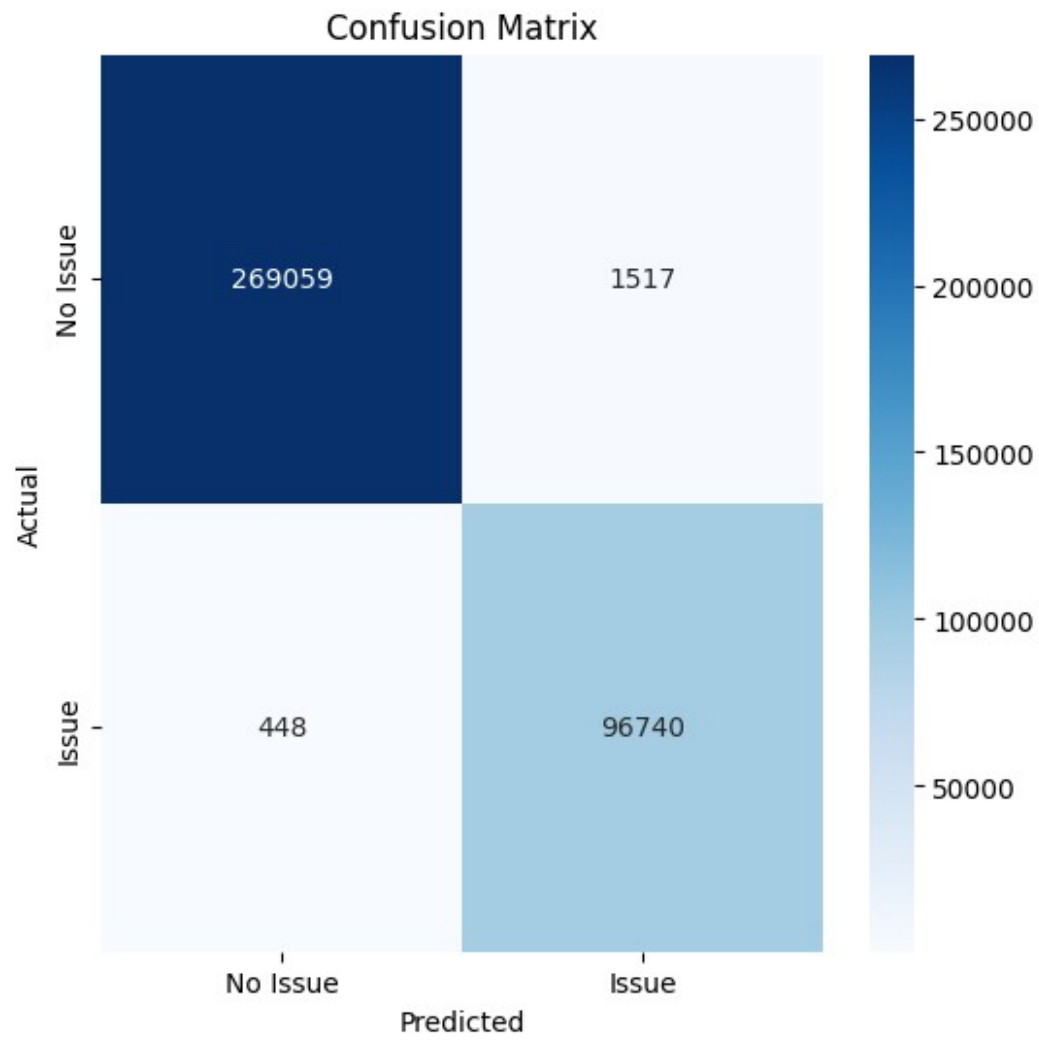
	precision	recall	f1-score	support
No Issue	1.00	0.99	1.00	8953
Issue	0.99	1.00	0.99	3306
accuracy			1.00	12259
macro avg	0.99	1.00	0.99	12259
weighted avg	1.00	1.00	1.00	12259


Epoch 3/6


Train Loss: 0.0000, Accuracy: 0.9976

Val Loss: 0.0221, Accuracy: 0.9967

	precision	recall	f1-score	support
No Issue	1.00	1.00	1.00	8953
Issue	0.99	1.00	0.99	3306
accuracy			1.00	12259
macro avg	0.99	1.00	1.00	12259







Average Metrics:

Average Accuracy: 0.9849

Macro-Precision: 0.9850

Macro-Recall: 0.9849

Macro-F1 Score: 0.9849

Weighted-Precision: 0.9850

Weighted-Recall: 0.9849

Weighted-F1 Score: 0.9849